

kNN & Naïve Bayes

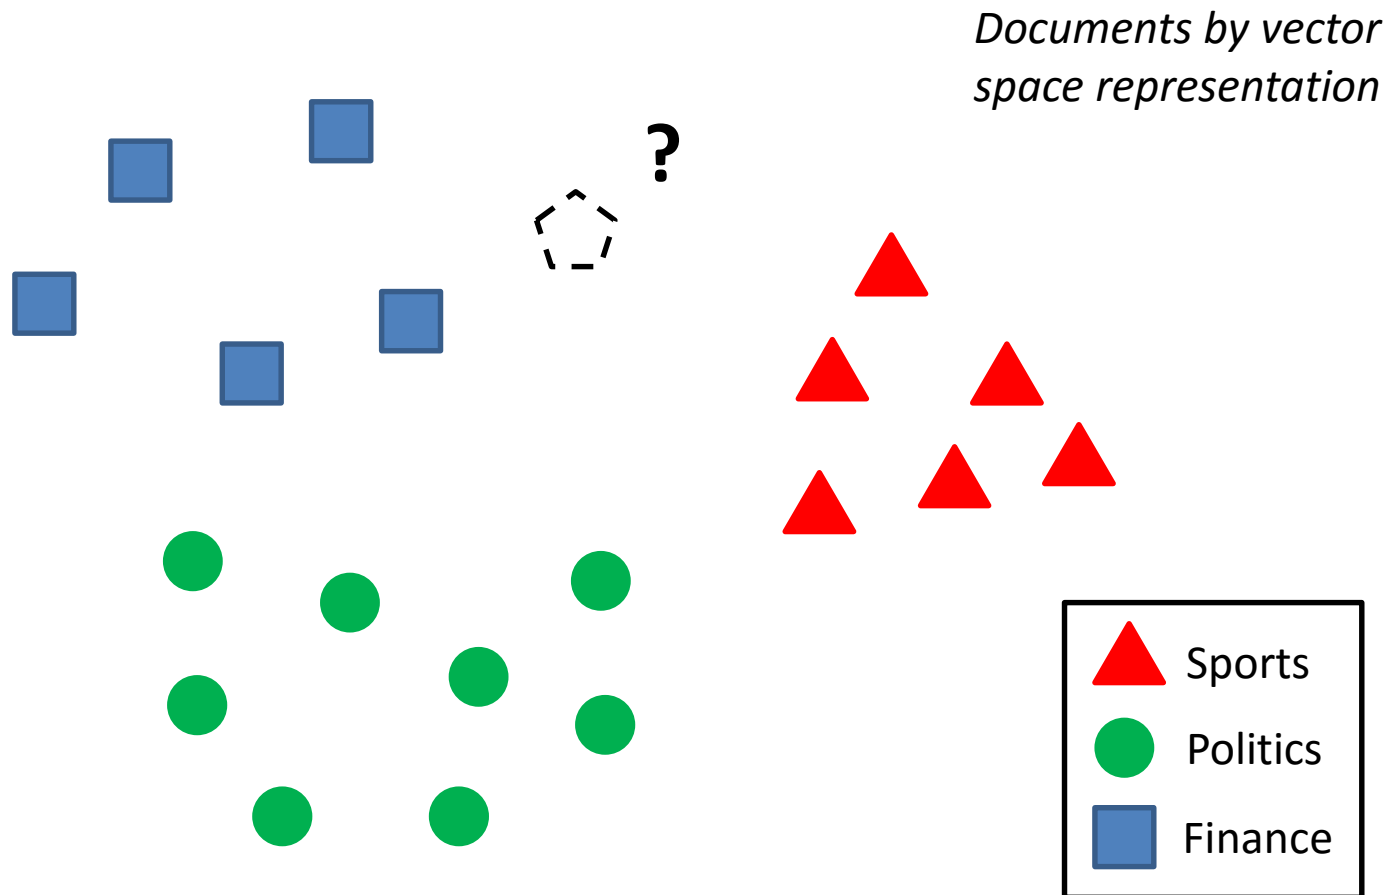
Hongning Wang

CS@UVa


Today's lecture

- Instance-based classifiers
 - k nearest neighbors
 - Non-parametric learning algorithm
- Model-based classifiers
 - Naïve Bayes classifier
 - A generative model
 - Parametric learning algorithm

How to classify this document?




Let's check the nearest neighbor



Let's check more nearest neighbors

- Ask k nearest neighbors

– Let them vote



Probabilistic interpretation of kNN

- Approximate Bayes decision rule in a subset of data around the testing point
- Let V be the volume of the m dimensional ball around x containing the k nearest neighbors for x , we have

$$p(x)V = \frac{k}{N} \Rightarrow p(x) = \frac{k}{NV}$$

Total number of instances

$$p(x|y=1) = \frac{k_1}{N_1V}$$

Nearest neighbors from class 1

$$p(y=1) = \frac{N_1}{N}$$

Total number of instances in class 1

With Bayes rule:


$$p(y=1|x) = \frac{\frac{N_1}{N} \times \frac{k_1}{N_1V}}{\frac{k}{NV}} = \frac{k_1}{k}$$

Counting the nearest neighbors from class 1

kNN is close to optimal

- Asymptotically, the error rate of 1-nearest-neighbor classification is less than twice of the Bayes error rate
- Decision boundary
 - 1NN - Voronoi tessellation

A non-parametric estimation of posterior distribution

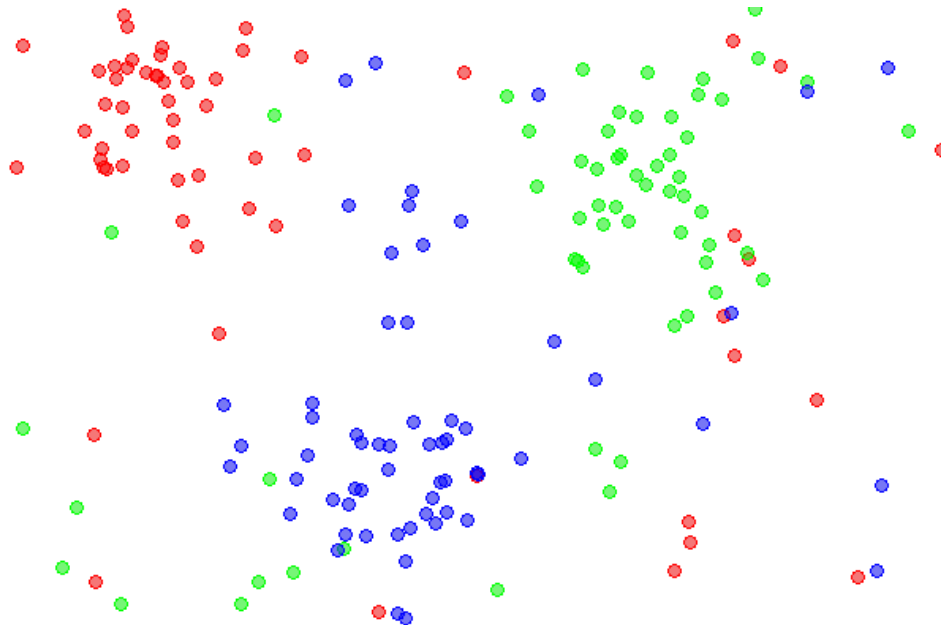


Components in kNN

- A distance metric
 - Euclidean distance/cosine similarity
- How many nearby neighbors to look at
 - k
- Instance look up
 - Efficiently search nearby points

Effect of k

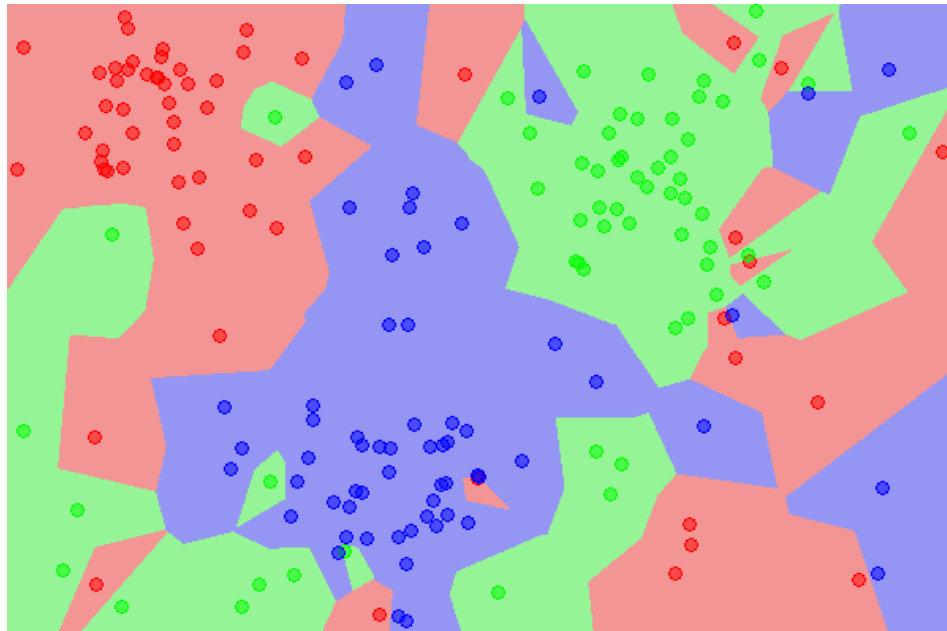
- Choice of k influences the “smoothness” of the resulting classifier



Effect of k

- Choice of k influences the “smoothness” of the resulting classifier

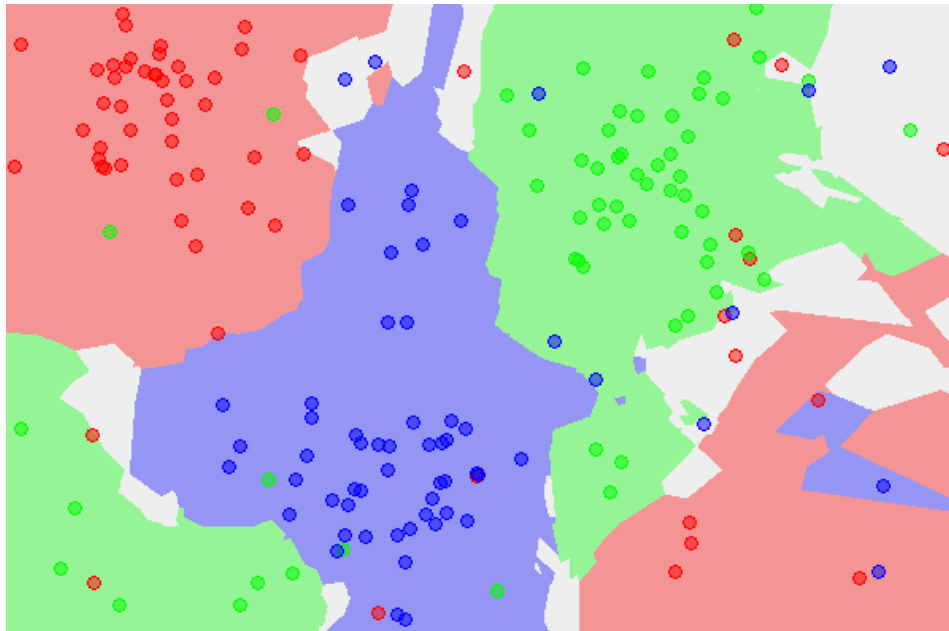
k=1



Effect of k


- Choice of k influences the “smoothness” of the resulting classifier

k=5

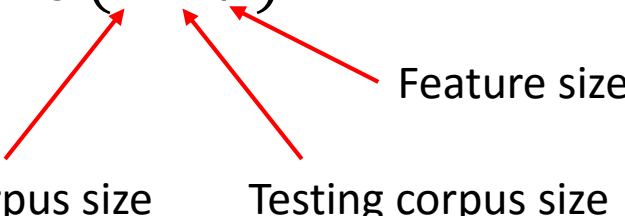


Effect of k

- Large k \rightarrow smooth shape for decision boundary
- Small k \rightarrow complicated decision boundary

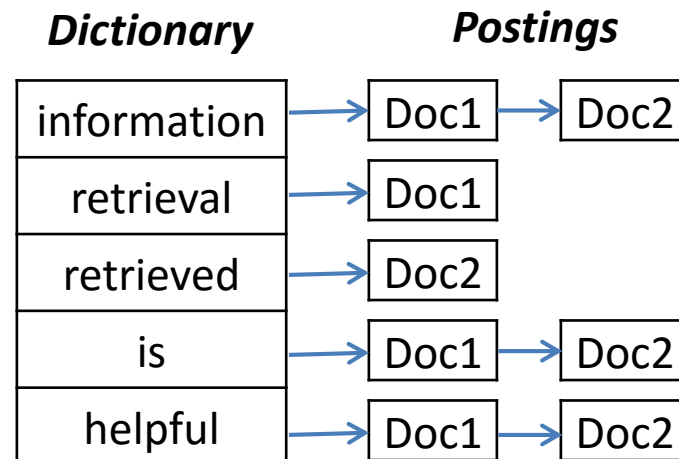


Efficient instance look-up

- Recall MP1
 - In Yelp_small data set, there are 629K reviews for training and 174K reviews for testing
 - Assume we have a vocabulary of 15k
 - Complexity of kNN
 - $O(NMV)$ 
 - Training corpus size
 - Testing corpus size
 - Feature size

Efficient instance look-up

- Exact solutions
 - Build inverted index for documents
 - Special mapping: word -> document list
 - Speed-up is limited when average document length is large




Efficient instance look-up

- Exact solutions
 - Build inverted index for documents
 - Special mapping: word -> document list
 - Speed-up is limited when average document length is large
 - Parallelize the computation
 - Map-Reduce
 - Map training/testing data onto different reducers
 - Merge the nearest k neighbors from the reducers

Efficient instance look-up

- Approximate solution
 - Locality sensitive hashing
 - Similar documents -> (likely) same hash values






Efficient instance look-up

- Approximate solution
 - Locality sensitive hashing
 - Similar documents -> (likely) same hash values
 - Construct the hash function such that similar items map to the same “buckets” with high probability
 - Learning-based: learn the hash function with annotated examples, e.g., must-link, cannot-link
 - Random projection

Random projection

- Approximate the cosine similarity between vectors
 - $h^r(x) = \text{sgn}(x \cdot r)$, r is a random unit vector
 - Each r defines one hash function, i.e., one bit in the hash value




	r_1	r_2	r_3
D_x	1	1	0
D_y	1	0	1



Random projection

- Approximate the cosine similarity between vectors
 - $h^r(x) = \text{sgn}(x \cdot r)$, r is a random unit vector
 - Each r defines one hash function, i.e., one bit in the hash value

	r_1	r_2	r_3
D_x	1	0	1
D_y	1	0	1




Random projection

- Approximate the cosine similarity between vectors
 - $h^r(x) = \text{sgn}(x \cdot r)$, r is a random unit vector
 - Each r defines one hash function, i.e., one bit in the hash value
 - Provable approximation error
 - $P(h(x) = h(y)) = 1 - \frac{\theta(x,y)}{\pi}$


Efficient instance look-up

- Effectiveness of random projection
 - 1.2M images + 1000 dimensions



Weight the nearby instances

- When the data distribution is highly skewed, frequent classes might dominate majority vote
 - They occur more often in the k nearest neighbors just because they have large volume



Weight the nearby instances


- When the data distribution is highly skewed, frequent classes might dominate majority vote
 - They occur more often in the k nearest neighbors just because they have large volume
- Solution
 - Weight the neighbors in voting
 - $w(x, x_i) = \frac{1}{|x - x_i|}$ or $w(x, x_i) = \cos(x, x_i)$

Summary of kNN

- Instance-based learning
 - No training phase
 - Assign label to a testing case by its nearest neighbors
 - Non-parametric
 - Approximate Bayes decision boundary in a local region
- Efficient computation
 - Locality sensitive hashing
 - Random projection

Recall optimal Bayes decision boundary

- $f(X) = \operatorname{argmax}_y P(y|X)$



Estimating the optimal classifier

- $$f(X) = \operatorname{argmax}_y P(y|X)$$

$$= \operatorname{argmax}_y P(X|y)P(y)$$

Requirement:

$$|D| \gg |Y| \times (2^V - 1)$$

Class conditional density

Class prior

#parameters:

$$|Y| \times (2^V - 1)$$

$$|Y| - 1$$

	text	information	identify	mining	mined	is	useful	to	from	apple	delicious	Y
D1	1	1	1	1	0	1	1	1	0	0	0	1
D2	1	1	0	0	1	1	1	0	1	0	0	1
D3	0	0	0	0	0	1	0	0	0	1	1	0

V binary features

We need to simplify this

- Features are conditionally independent given class label
 - $p(x_1, x_2|y) = p(x_2|x_1, y)p(x_1|y)$
 $= p(x_2|y)p(x_1|y)$
 - E.g.,
 $p(\text{'white house'}, \text{'obama'} | \text{political news}) =$
 $p(\text{'white house'} | \text{political news}) \times$
 $p(\text{'obama'} | \text{political news})$

This does not mean 'white house' is independent of 'obama'!

Conditional v.s. marginal independence

- Features are not necessarily marginally independent from each other
 - $p(\text{'white house'} | \text{'obama'}) > p(\text{'white house'})$
- However, once we know the class label, features become independent from each other
 - Knowing it is already political news, observing 'obama' contributes little about occurrence of 'while house'

Naïve Bayes classifier

- $f(X) = \operatorname{argmax}_y P(y|X)$
 $= \operatorname{argmax}_y P(X|y)P(y)$

$$= \operatorname{argmax}_y \prod_{i=1}^{|d|} P(x_i|y) P(y)$$

Class conditional density

Class prior

#parameters:

$$|Y| \times V$$

$$|Y| - 1$$


v.s.

$$|Y| \times (2^V - 1)$$

Computationally feasible

Naïve Bayes classifier

- $f(X) = \operatorname{argmax}_y P(y|X)$
 $= \operatorname{argmax}_y P(X|y)P(y)$ *By Bayes rule*
 $= \operatorname{argmax}_y \prod_{i=1}^{|d|} P(x_i|y) P(y)$ *By independence assumption*



Estimating parameters

- Maximial likelihood estimator

$$- P(x_i|y) = \frac{\sum_d \sum_j \delta(x_d^j = x_i, y_d = y)}{\sum_d \delta(y_d = y)}$$

$$- P(y) = \frac{\sum_d \delta(y_d = y)}{\sum_d 1}$$

	text	information	identify	mining	mined	is	useful	to	from	apple	delicious	Y
D1	1	1	1	1	0	1	1	1	0	0	0	1
D2	1	1	0	0	1	1	1	0	1	0	0	1
D3	0	0	0	0	0	1	0	0	0	1	1	0

Enhancing Naïve Bayes for text classification I

- The frequency of words in a document matters

- $P(X|y) = \prod_{i=1}^{|d|} P(x_i|y)^{c(x_i,d)}$

- In log space *Essentially, estimating $|Y|$ different language models!*

- $f(X) = \operatorname{argmax}_y \log P(y|X)$
 $= \operatorname{argmax}_y \log P(y) + \sum_{i=1}^{|d|} c(x_i, d) \log P(x_i|y)$

Class bias

Feature vector

Model parameter

Enhancing Naïve Bayes for text classification

- For binary case

$$\begin{aligned} -f(X) &= \text{sgn} \left(\log \frac{P(y = 1|X)}{P(y = 0|X)} \right) \\ &= \text{sgn} \left(\log \frac{P(y = 1)}{P(y = 0)} + \sum_{i=1}^{|d|} c(x_i, d) \log \frac{P(x_i|y = 1)}{P(x_i|y = 0)} \right) \\ &= \text{sgn}(w^T \bar{x}) \end{aligned}$$

a linear model with vector space representation?

where

$$\begin{aligned} w &= \left(\log \frac{P(y = 1)}{P(y = 0)}, \log \frac{P(x_1|y = 1)}{P(x_1|y = 0)}, \dots, \log \frac{P(x_v|y = 1)}{P(x_v|y = 0)} \right) \\ \bar{x} &= (1, c(x_1, d), \dots, c(x_v, d)) \end{aligned}$$

We will come back to this topic later.

Enhancing Naïve Bayes for text classification II

- Usually, features are not conditionally independent
 - $p(X|y) \neq \prod_{i=1}^{|d|} P(x_i|y)$
- Enhance the conditional independence assumptions by N-gram language models
 - $p(X|y) = \prod_{i=1}^{|d|} P(x_i|x_{i-1}, \dots, x_{i-N+1}, y)$

Enhancing Naïve Bayes for text classification III

- Sparse observation
 - $\delta(x_d^j = x_i, y_d = y) = 0 \Rightarrow p(x_i|y) = 0$
 - Then, no matter what values the other features take, $p(x_1, \dots, x_i, \dots, x_V|y) = 0$
- Smoothing class conditional density
 - All smoothing techniques we have discussed in language models are applicable here

Maximum a Posterior estimator

- Adding pseudo instances

- Priors: $q(y)$ and $q(x, y)$

Can be estimated from a related corpus or manually tuned

- MAP estimator for Naïve Bayes

- $$P(x_i|y) = \frac{\sum_d \sum_j \delta(x_d^j = x_i, y_d = y) + Mq(x_i, y)}{\sum_d \delta(y_d = y) + Mq(y)}$$

#pseudo instances

Summary of Naïve Bayes

- Optimal Bayes classifier
 - Naïve Bayes with independence assumptions
- Parameter estimation in Naïve Bayes
 - Maximum likelihood estimator
 - Smoothing is necessary

Today's reading

- Introduction to Information Retrieval
 - Chapter 13: Text classification and Naive Bayes
 - 13.2 – Naive Bayes text classification
 - 13.4 – Properties of Naive Bayes
 - Chapter 14: Vector space classification
 - 14.3 k nearest neighbor
 - 14.4 Linear versus nonlinear classifiers