

Part-of-Speech Tagging & Sequence Labeling

Hongning Wang

CS@UVa

What is POS tagging

Tag Set

NNP: proper noun

CD: numeral

JJ: adjective

POS Tagger

Raw Text

Pierre Vinken , 61 years
old, will join the board as
a nonexecutive director
Nov. 29 .

Tagged Text

Pierre_**NNP** Vinken_**NNP** ,_
61_**CD** years_**NNS** old_**JJ** ,_
will_**MD** join_**VB** the_**DT**
board_**NN** as_**IN** a_**DT**
nonexecutive_**JJ** director_**NN**
Nov._**NNP** 29_**CD** ._
.

Why POS tagging?

- POS tagging is a prerequisite for further NLP analysis
 - Syntax parsing
 - Basic unit for parsing
 - Information extraction
 - Indication of names, relations
 - Machine translation
 - The meaning of a particular word depends on its POS tag
 - Sentiment analysis
 - Adjectives are the major opinion holders
 - Good v.s. Bad, Excellent v.s. Terrible

Challenges in POS tagging

- Words often have more than one POS tag
 - The back door (adjective)
 - On my back (noun)
 - Promised to back the bill (verb)
- Simple solution with dictionary look-up does not work in practice
 - One needs to determine the POS tag for a particular instance of a word from its context

Define a tagset

- We have to agree on a standard inventory of word classes
 - Taggers are trained on a labeled corpora
 - The tagset needs to capture semantically or syntactically important distinctions that can easily be made by trained human annotators

Word classes

- Open classes
 - Nouns, verbs, adjectives, adverbs
- Closed classes
 - Auxiliaries and modal verbs
 - Prepositions, Conjunctions
 - Pronouns, Determiners
 - Particles, Numerals

Public tagsets in NLP

- Brown corpus - Francis and Kucera 1961
 - 500 samples, distributed across 15 genres in rough proportion to the amount published in 1961 in each of those genres
 - 87 tags
- Penn Treebank - Marcus et al. 1993
 - Hand-annotated corpus of Wall Street Journal, 1M words
 - 45 tags, a simplified version of Brown tag set
 - Standard for English now
 - Most statistical POS taggers are trained on this Tagset

How much ambiguity is there?

- Statistics of word-tag pair in Brown Corpus and Penn Treebank

		87-tag Original Brown		45-tag Treebank Brown	
Unambiguous (1 tag)		44,019		38,857	
Ambiguous (2–7 tags)		5,490	11%	8844	18%
Details:	2 tags	4,967		6,731	
	3 tags	411		1621	
	4 tags	91		357	
	5 tags	17		90	
	6 tags	2	(<i>well, beat</i>)	32	
	7 tags	2	(<i>still, down</i>)	6	(<i>well, set, round, open, fit, down</i>)
	8 tags			4	(<i>'s, half, back, a</i>)
	9 tags			3	(<i>that, more, in</i>)

Is POS tagging a solved problem?

- Baseline
 - Tag every word with its most frequent tag
 - Tag unknown words as nouns
 - Accuracy
 - Word level: 90%
 - Sentence level
 - Average English sentence length 14.3 words
 - $0.9^{14.3} = 22\%$
- Accuracy of State-of-the-art POS Tagger*
 - *Word level: 97%*
 - *Sentence level: $0.97^{14.3} = 65\%$*

Building a POS tagger

- Rule-based solution
 1. Take a dictionary that lists all possible tags for each word
 2. Assign to every word all its possible tags
 3. Apply rules that eliminate impossible/unlikely tag sequences, leaving only one tag per word


she PRP
promised ~~VDN,VBD~~
to TO
back VB, JJ, RB, NN!!
the DT
bill NN, ~~VB~~

*R1: Pronoun should be
followed by a past tense verb*

*R2: Verb cannot follow
determiner*

Building a POS tagger

- Statistical POS tagging



- What is the most likely sequence of tags t for the given sequence of words w

$$t^* = \operatorname{argmax}_t p(t|w)$$

POS tagging with generative models

- Bayes Rule

$$\begin{aligned} \mathbf{t}^* &= \operatorname{argmax}_{\mathbf{t}} p(\mathbf{t}|\mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{t}} \frac{p(\mathbf{w}|\mathbf{t})p(\mathbf{t})}{p(\mathbf{w})} \\ &= \operatorname{argmax}_{\mathbf{t}} \boxed{p(\mathbf{w}|\mathbf{t})} \boxed{p(\mathbf{t})} \end{aligned}$$

- Joint distribution of tags and words

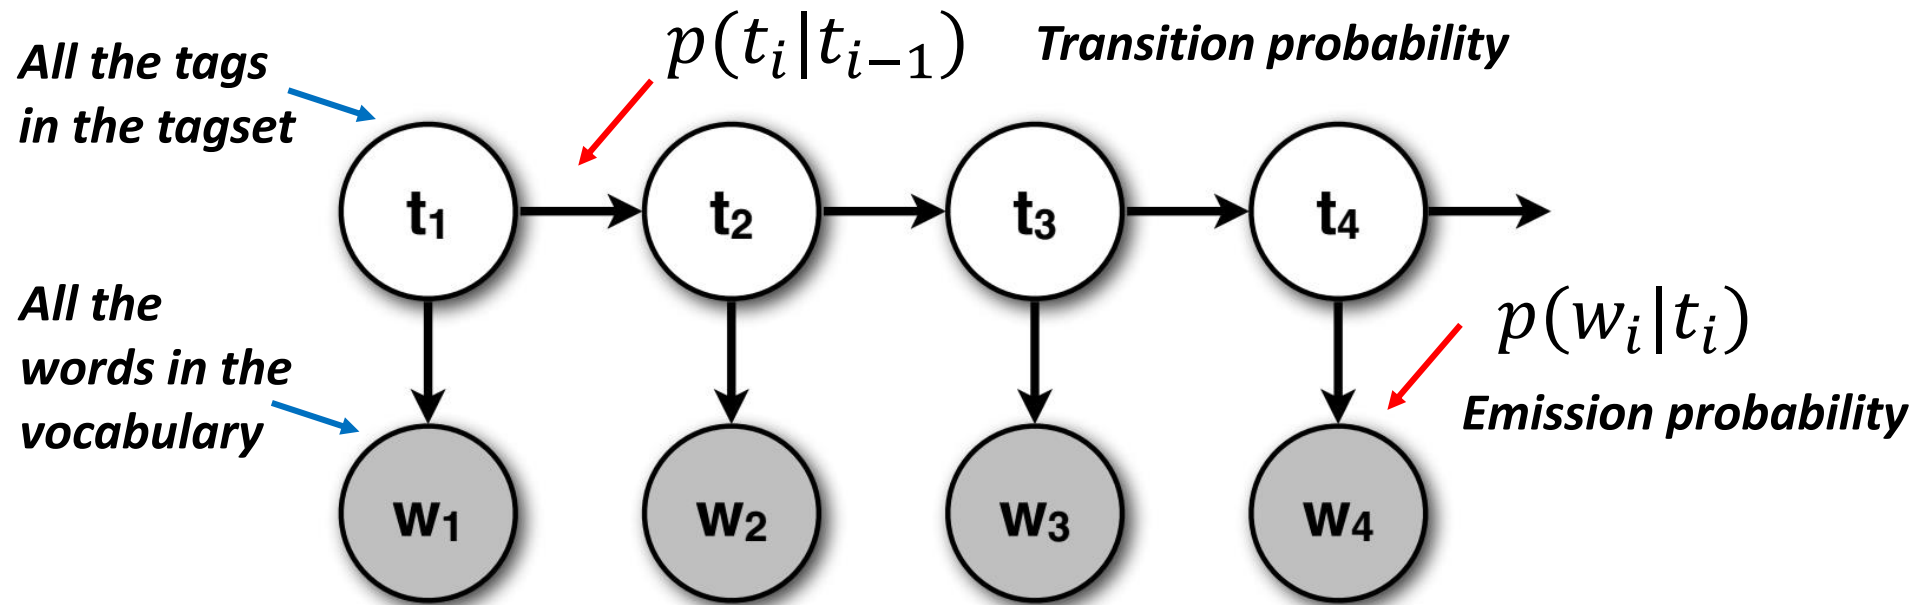
- Generative model

- A stochastic process that **first generates the tags**, and then **generates the words based on these tags**

Hidden Markov models

- Two assumptions for POS tagging
 1. Current tag only depends on previous k tags
 - $p(\mathbf{t}) = \prod_i p(t_i | t_{i-1}, t_{i-2}, \dots, t_{i-k})$
 - When $k=1$, it is so-called first-order HMMs
 2. Each word in the sequence depends only on its corresponding tag
 - $p(\mathbf{w} | \mathbf{t}) = \prod_i p(w_i | t_i)$

Graphical representation of HMMs



- Light circle: latent random variables
- Dark circle: observed random variables
- Arrow: probabilistic dependency

Finding the most probable tag sequence


$$\mathbf{t}^* = \operatorname{argmax}_{\mathbf{t}} p(\mathbf{t}|\mathbf{w})$$

$$= \operatorname{argmax}_{\mathbf{t}} \prod_i p(w_i|t_i)p(t_i|t_{i-1})$$

- Complexity analysis
 - Each word can have up to T tags
 - For a sentence with N words, there will be up to T^N possible tag sequences
 - Key: explore the **special structure** in HMMs!

Trellis: a special structure for HMMs

	w_1	w_2	w_3	w_4	w_5
t_1					
t_2					
t_3					
t_4					
t_5					
t_6					
t_7					




Word w_1 takes tag t_4 .

Trellis: a special structure for HMMs

$$t^1 = \underline{t_4} \underline{t_1} \underline{t_3} \underline{t_5} \boxed{t_7}$$

$$t^2 = \underline{t_4} \underline{t_1} \underline{t_3} \underline{t_5} \boxed{t_2}$$

Computation can be reused!



Word w_1 takes tag t_4 .

Viterbi algorithm

- Store the best tag sequence for $w_1 \dots w_i$ that ends in t^j in $T[j][i]$

$$- T[j][i] = \max p(w_1 \dots w_i, t_1 \dots, t_i = t^j)$$

- Recursively compute $\text{trellis}[j][i]$ from the entries in the previous column $\text{trellis}[j][i-1]$

$$- T[j][i] = P(w_i | t^j) \text{Max}_k \left(T[k][i-1] P(t^j | t_k) \right)$$

Generating the current observation

The best i-1 tag sequence

Transition from the previous best ending tag

Dynamic programming: $O(TN)$!

Decode $\operatorname{argmax}_t p(t, w)$

- Take the highest scoring entry in the last column of the trellis

Keep backpointers in each trellis to keep track of the most probable sequence

	w_1	w_2	w_3	w_4	w_5
t_1					
t_2					
t_3					
t_4					
t_5					
t_6					
t_7					

Train an HMMs tagger

- Parameters in an HMMs tagger
 - Transition probability: $p(t_i|t_j), T \times T$
 - Emission probability: $p(w|t), V \times T$
 - Initial state probability: $p(t|\pi), T \times 1$



For the first tag in a sentence

Train an HMMs tagger

- Maximum likelihood estimator
 - Given a labeled corpus, e.g., Penn Treebank
 - Count how often we have the pair of $t_i t_j$ and $w_i t_j$
 - $p(t_i | t_j) = \frac{c(t_i, t_j)}{c(t_j)}$
 - $p(w_i | t_j) = \frac{c(w_i, t_j)}{c(t_j)}$

Public POS taggers

- Brill's tagger
 - <http://www.cs.jhu.edu/~brill/>
- TnT tagger
 - <http://www.coli.uni-saarland.de/~thorsten/tnt/>
- Stanford tagger
 - <http://nlp.stanford.edu/software/tagger.shtml>
- SVMTool
 - <http://www.lsi.upc.es/~nlp/SVMTool/>
- GENIA tagger
 - <http://www-tsujii.is.s.u-tokyo.ac.jp/GENIA/tagger/>
- More complete list at
 - <http://www-nlp.stanford.edu/links/statnlp.html#Taggers>

Let's take a look at other NLP tasks

- Noun phrase (NP) chunking
 - Task: identify all non-recursive NP chunks

Pierre Vinken , 61 years old , will join IBM 's board
as a nonexecutive director Nov. 29 .



[NP Pierre Vinken] , [NP 61 years] old , will join
[NP IBM] 's [NP board] as [NP a nonexecutive director]
[NP Nov. 2] .

The BIO encoding

- Define three new tags
 - B-NP: beginning of a noun phrase chunk
 - I-NP: inside of a noun phrase chunk
 - O: outside of a noun phrase chunk

[NP Pierre Vinken] , [NP 61 years] old , will join
[NP IBM] 's [NP board] as [NP a nonexecutive director]
[NP Nov. 2] .



POS Tagging with restricted Tagset?

Pierre_B-NP Vinken_I-NP ,_O 61_B-NP years_I-NP
old_O ,_O will_O join_O IBM_B-NP 's_O board_B-NP as_O
a_B-NP nonexecutive_I-NP director_I-NP Nov._B-NP
29_I-NP ._O

Another NLP task

- Shallow parsing
 - Task: identify all non-recursive NP, verb (“VP”) and preposition (“PP”) chunks

Pierre Vinken , 61 years old , will join IBM 's board
as a nonexecutive director Nov. 29 .



[NP Pierre Vinken] , [NP 61 years] old , [VP will join]
[NP IBM] 's [NP board] [PP as] [NP a nonexecutive
director] [NP Nov. 2] .

BIO Encoding for Shallow Parsing

- Define several new tags
 - B-NP B-VP B-PP: beginning of an NP, “VP”, “PP” chunk
 - I-NP: inside of an NP, “VP”, “PP” chunk
 - O: outside of any chunk

[NP Pierre Vinken] , [NP 61 years] old , [VP will join]
[NP IBM] 's [NP board] [PP as] [NP a nonexecutive
director] [NP Nov. 2] .



POS Tagging with restricted Tagset?

Pierre_B-NP Vinken_I-NP ,_O 61_B-NP years_I-NP
old_O ,_O will_B-VP join_I-VP IBM_B-NP 's_O board_B-NP
as_B-PP a_B-NP nonexecutive_I-NP director_I-NP Nov._B-
NP 29_I-NP ._O

Yet Another NLP task

- Named Entity Recognition
 - Task: identify all mentions of named entities (people, organizations, locations, dates)

Pierre Vinken , 61 years old , will join IBM 's board
as a nonexecutive director Nov. 29 .



[PERS Pierre Vinken] , 61 years old , will join
[ORG IBM] 's board as a nonexecutive director
[DATE Nov. 2] .

BIO Encoding for NER

- Define many new tags
 - B-PERS, B-DATE,...: beginning of a mention of a person/date...
 - I-PERS, B-DATE,...: inside of a mention of a person/date...
 - O: outside of any mention of a named entity

[PERS Pierre Vinken] , 61 years old , will join
[ORG IBM] 's board as a nonexecutive director
[DATE Nov. 2] .



POS Tagging with restricted Tagset?

Pierre_B-PERS Vinken_I-PERS ,_O 61_O years_O old_O ,_O
will_O join_O IBM_B-ORG 's_O board_O as_O a_O
nonexecutive_O director_O Nov._B-DATE 29_I-DATE ._O

Sequence labeling

- Many NLP tasks are sequence labeling tasks
 - Input: a sequence of tokens/words
 - Output: a sequence of corresponding labels
 - E.g., POS tags, BIO encoding for NER
 - Solution: finding the most probable label sequence for the given word sequence
 - $\mathbf{t}^* = \operatorname{argmax}_{\mathbf{t}} p(\mathbf{t}|\mathbf{w})$

Comparing to traditional classification problem

Sequence labeling

- $\mathbf{t}^* = \operatorname{argmax}_{\mathbf{t}} p(\mathbf{t}|\mathbf{w})$
 - \mathbf{t} is a vector/matrix
- Dependency between both (\mathbf{t}, \mathbf{w}) and (\mathbf{t}, \mathbf{t})
- Structured output
- Difficult to solve the inference problem

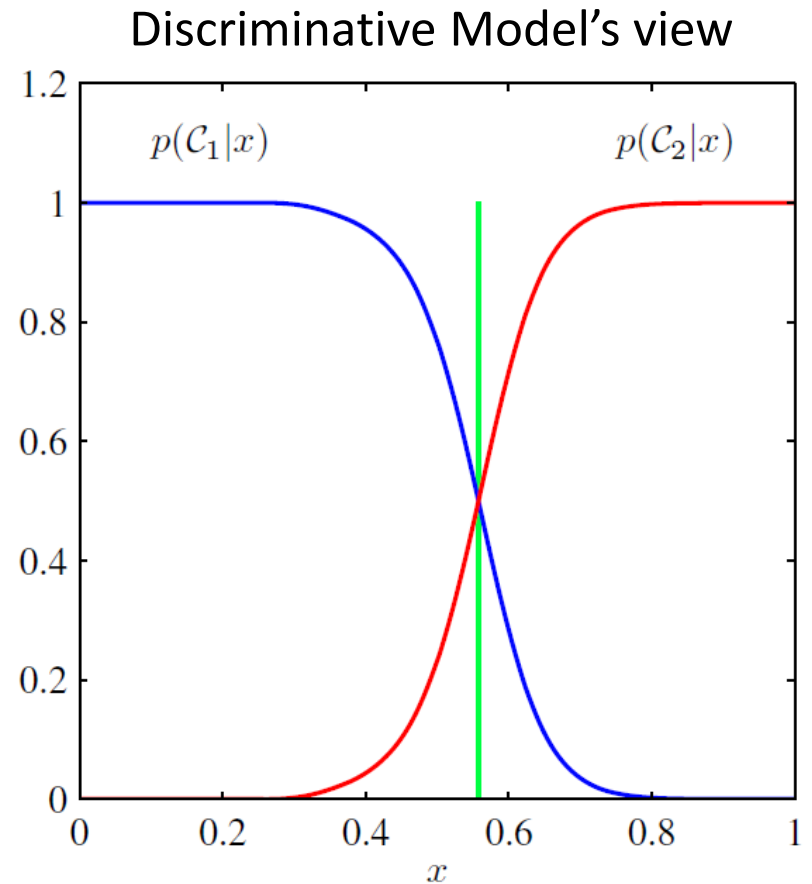

Traditional classification

- $y = \operatorname{argmax}_y p(y|\mathbf{x})$
 - y is a single label
- Dependency only within (\mathbf{t}, \mathbf{w})
- Independent output
- Easy to solve the inference problem

Two modeling perspectives

- Generative models
 - Model the joint probability of labels and words
 - $\mathbf{t}^* = \operatorname{argmax}_{\mathbf{t}} p(\mathbf{t}|\mathbf{w}) = \operatorname{argmax}_{\mathbf{t}} p(\mathbf{w}|\mathbf{t})p(\mathbf{t})$
- Discriminative models
 - Directly model the conditional probability of labels given the words
 - $\mathbf{t}^* = \operatorname{argmax}_{\mathbf{t}} p(\mathbf{t}|\mathbf{w}) = f(\mathbf{t}, \mathbf{w})$

Generative V.S. discriminative models



Generative V.S. discriminative models

Generative


- Specifying joint distribution
 - Full probabilistic specification for all the random variables
- Dependence assumption has to be specified for $p(\mathbf{w}|\mathbf{t})$ and $p(\mathbf{t})$
- Flexible, can be used in unsupervised learning

Discriminative

- Specifying conditional distribution
 - Only explain the target variable
- Arbitrary features can be incorporated for modeling $p(\mathbf{t}|\mathbf{w})$
- Need training data, only suitable for (semi-) supervised learning

Maximum entropy Markov models

- MEMMs are discriminative models of the labels \mathbf{t} given the observed input sequence \mathbf{w}
 - $p(\mathbf{t}|\mathbf{w}) = \prod_i p(t_i|w_i, t_{i-1})$



Design features

- Emission-like features

- Binary feature functions

- $f_{\text{first-letter-capitalized-NNP}}(\text{China}) = 1$
 - $f_{\text{first-letter-capitalized-VB}}(\text{fly}) = 0$


- Integer (or real-valued) feature functions

- $f_{\text{number-of-vowels-NNP}}(\text{China}) = 2$

- Transition-like features

- Binary feature functions

- $f_{\text{first-letter-capitalized-NNP-VB}}(\text{China}) = 1$




Not necessarily
independent features!

Parameterization of $p(t_i | w_i, t_{i-1})$

- Associate a real-valued weight λ to each specific type of feature function
 - λ_k for $f_{\text{first-letter-capitalized-NNP}}(w)$
- Define a scoring function $f(t_i, t_{i-1}, w_i) = \sum_k \lambda_k f_k(w_i)$
- Naturally $p(t_i | w_i, t_{i-1}) \propto \exp f(t_i, t_{i-1}, w_i)$
 - Recall the basic definition of probability
 - $P(x) > 0$
 - $\sum_x p(x) = 1$

Parameterization of MEMMs


$$p(\mathbf{t}|\mathbf{w}) = \prod_i p(t_i|w_i, t_{i-1})$$
$$\propto \prod_i \exp(f(t_i, t_{i-1}, w_i))$$

- It is a log-linear model
– $-\log p(\mathbf{t}|\mathbf{w}) = \sum_i f(t_i, t_{i-1}, w_i) - \mathcal{C}(\lambda)$
Constant only related to λ 
- Viterbi algorithm can be used to decode the most probable label sequence solely based on $\sum_i f(t_i, t_{i-1}, w_i)$

Parameter estimation


- Maximum likelihood estimator can be used in a similar way as in HMMs

$$\begin{aligned} -\lambda^* &= \operatorname{argmax}_{\lambda} \sum_{t,w} \log p(t|w) \\ &= \operatorname{argmax}_{\lambda} \sum_{t,w} \sum_i f(t_i, t_{i-1}, w_i) - C(\lambda) \end{aligned}$$




Why maximum entropy?

- We will explain this in detail when discussing the Logistic Regression models



A little bit more about MEMMs

- Emission features can go across multiple observations
 - $f(t_i, t_{i-1}, w_i) = \sum_k \lambda_k f_k(w_i, \mathbf{w})$
 - Especially useful for shallow parsing and NER tasks



Conditional random field

- A more advanced model for sequence labeling
 - Model global dependency
 - $p(t|w) \propto \prod_i \exp(\sum_k \lambda_k f_k(t_i, \mathbf{w}) + \sum_l \eta_l g_l(t_i, t_{i-1}, \mathbf{w}))$

